

The Google™ File System

King.Zevin@ADSLab



Outline

- Motivation
- Assumptions
- Questions
- Measurements
- Benefits/Limitations
- Conclusion

Motivation

- Need for a scalable DFS
- Large distributed data-intensive applications
- High data processing needs
- Performance, Reliability, Scalability and Availability
- More than traditional DFS

Assumptions

- Assumptions – Environment
- Assumptions – Applications

Assumptions – Environment

- Commodity Hardware
 - inexpensive
- TBs of Space
 - must support TBs of space
- Component Failure
 - the norm rather than the exception


Assumptions – Applications

- **Multi-GB files**
 - Common
- **Workloads**
 - Large streaming reads
 - Small random reads
 - Large, sequential writes that append data to file
 - Multiple clients concurrently append to one file
- **High sustained bandwidth**
 - More important than low latency

Question 0 – How to design – isolation

- **large data and small control infos**
 - Separate metadata from data
- **Then we get the initial architecture model**
 - Client – where user code runs
 - Master – metadata for the chunks
 - Chunkservers – large data

Architecture model



Question 0 – How to design – isolation

- **large data and small control infos**
 - Separate metadata from data
- **Then we get the architecture model**
 - Client – where user code runs
 - Master – metadata for the chunks
 - Chunkservers – large data
- **Then some questions will be brought up**

More questions – single master

- **What if the single master fails?**
 - Single master is not single – we make replicas
- **Network bottleneck?**
 - Reduce the interaction with master
 - Let the clients hold sth to interact with chunkservers more
 - Cache – clients cache the chunk infos
 - Let the chunkservers can interact more with each other
 - Primary and secondary
- **Too much metadata to hold**
- **No disks! All in memory!**

Question 1 – Too much metadata

- Too much metadata to hold and we need
- No disks! All in memory!
- Traditional block is too small
- Let the chunk size be 64MB

Benefits – chunk size : 64MB

- **Larger chunk, less metadata**
 - Master maintains less than 64B of metadata for 64MB chunk
 - 1 TB chunks : < 1 MB metadata
 - 1 PB chunks : < 1 GB metadata
- **Less interactions between client and master**
- **One continuous TCP connection during one chunk**
- **One disadvantage**
 - little files with one chunk
might make the chunkserver hotspot

So we get “chunk”

- Files are divided into *chunks*
- Replicated over *chunkservers*, called *replicas*
- Unique 64-bit *chunk handles*
- Chunks as Linux files

Question 2 – What to hold in master

- **Metadata :**
 - The file and chunk namespaces
 - Mapping from files to chunks
 - Locations of each chunk's replicas
- **All in memory**
- **first two kept persistent by operation log**
- **Get locations when startup and with heartbeat**
 - Change too often

Question 2.1 – Why operation log?

- We need reliability and ability to recovery
- Master might need restart, might fail
- Critical historical logs must be kept
 - In disks
 - By multiple servers
- In addition to the logs, we need to
- Checkpoint
- When checkpointing, create a new operation log and start a new checkpointing thread

Question 3 – What master does

- **Chunk creation**
- **Re-replication**
- **Garbage collection**

Question 3 – What master does

- **Chunk creation**
 - Low disk usage
 - Limit recent creations
 - Different racks
- **Re-replication**
- **Garbage collection**

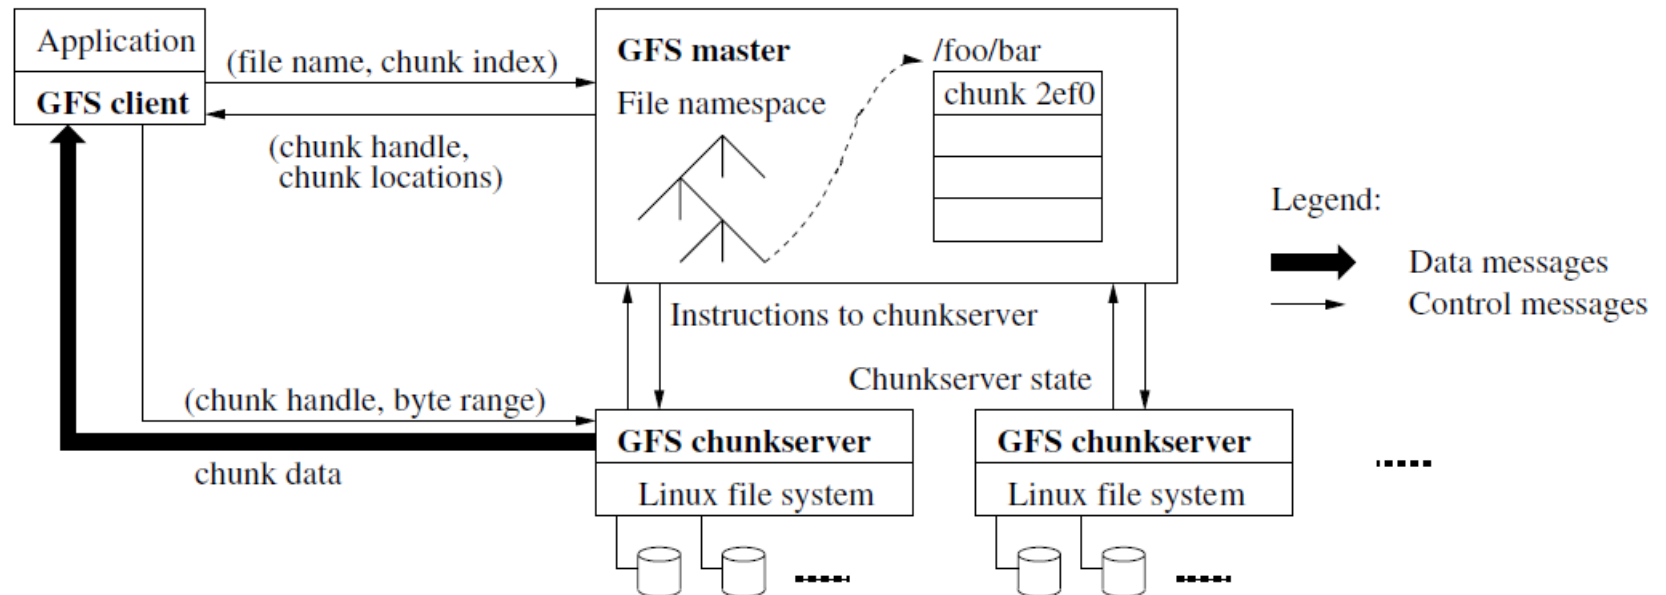
Question 3 – What master does

- **Chunk creation**
- **Re-replication**
 - Priority
 - 2 failure > 1 failure
 - Live > deleted
 - Blocking > not blocking
- **Garbage collection**

Question 3 – What master does

- **Chunk creation**
- **Re-replication**
- **Garbage collection**
 - Does not delete immediately unless deleted twice
 - Why?
 - Simple and amortized(background activity, done only when free)
 - How?
 - After deleted, file is renamed to a hidden name
 - Deleted after 3 days
 - Can be undeleted and read

Architecture model



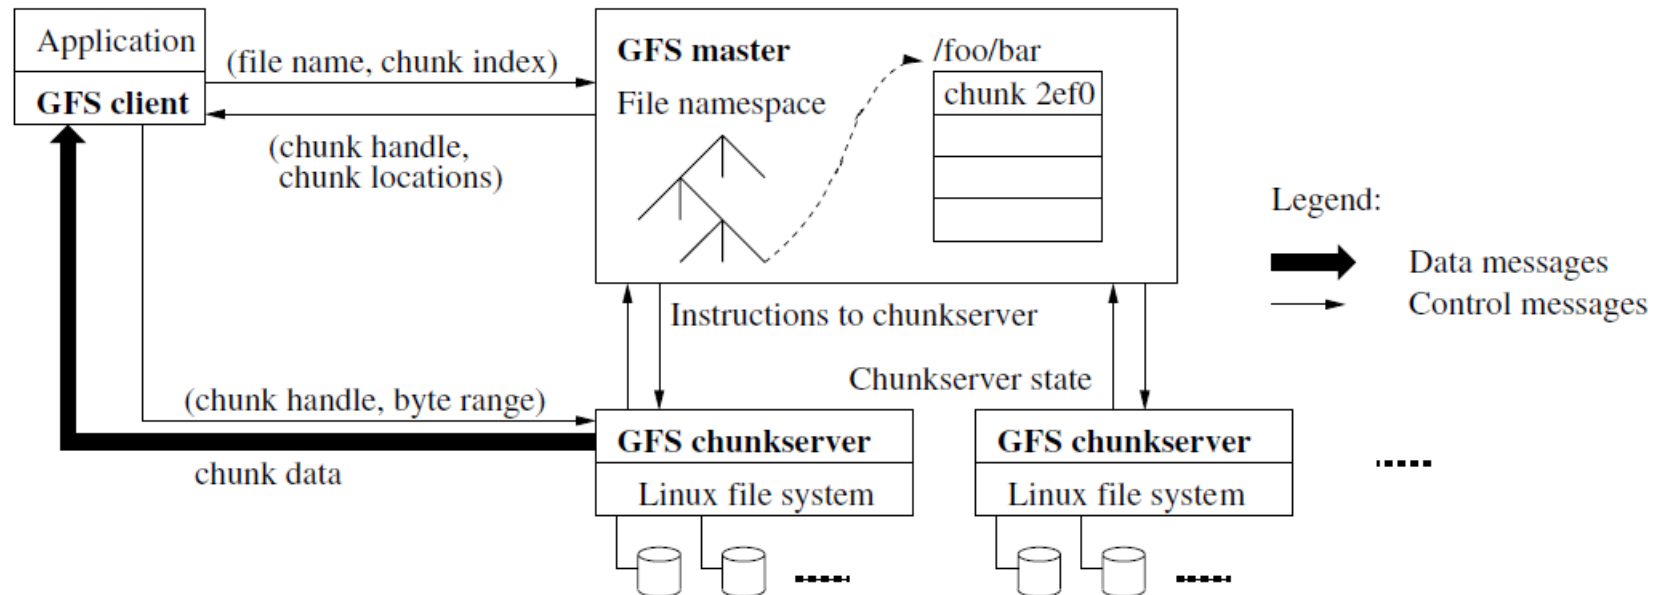
Question 4 – Namespace Management and Locking

- We need concurrency, so we need locking
- No per-directory data structure.
- No hard or symbolic links
- Namespaces are represented as a lookup table mapping full pathnames to metadata
 - Prefix compression
- Each master operation acquires a set of locks before it runs

Question 4 – Example of Locking Mechanism

- Preventing /home/user/foo from being created while /home/user is being snapshotted to /save/user
 - Snapshot operation
 - Read locks on /home and /save
 - Write locks on /home/user and /save/user
 - File creation
 - read locks on /home and /home/user
 - write locks on /home/user/foo
 - Conflict locks on /home/user

Architecture model



Question 5 – How to minimize master interactions


- Master – too much involvement!
- Between master and client, we use cache.
- Then what about master and chunkservers?
- Then give more power to chunkservers!
- Lease!

Question 5 –

Lease: to minimize master interactions

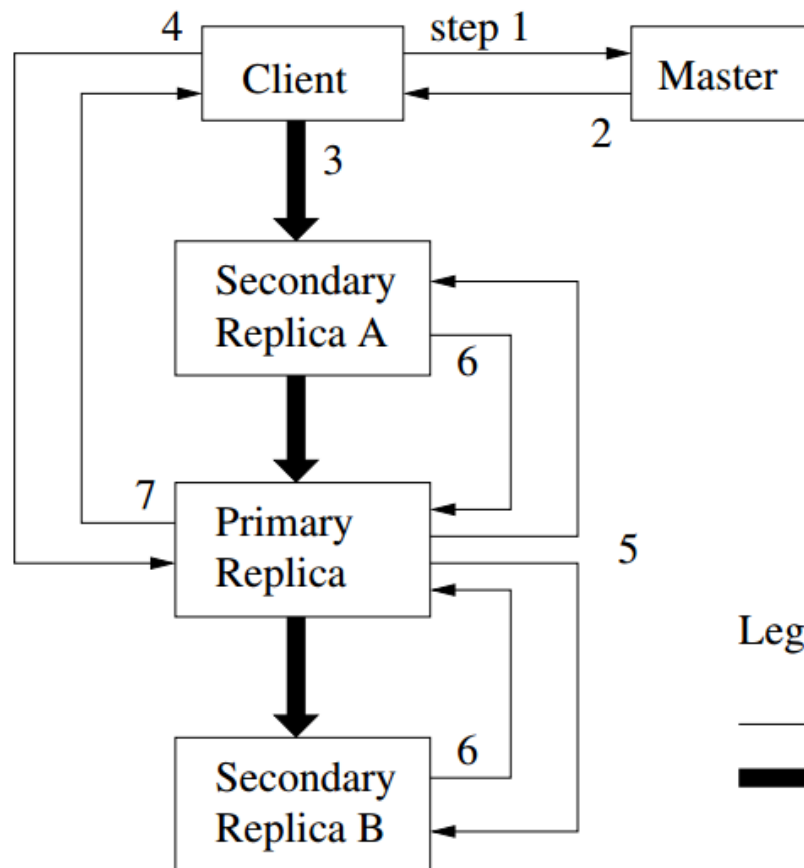
- Master grants a lease to a chunk – primary chunk
 - Primary chunk picks the mutation order, which secondary chunks will follow.
 - A lease is about 60s, extension requests can be piggybacked on HeartBeat message.
-
- Then we get the more detailed picture for the architecture model.

Question 6 – How the components interact



Step1:
Asks which
chunkserver
holds lease
and locations
of all replicas.
If no lease,
grant one.

Question 6 – How the components interact




Step2:
Master replies.
Client caches
locations,
contact master
again only when
prim
unreachable or
prim is not prim.

Legend:

→ Control
→ Data

Question 6 – How the components interact

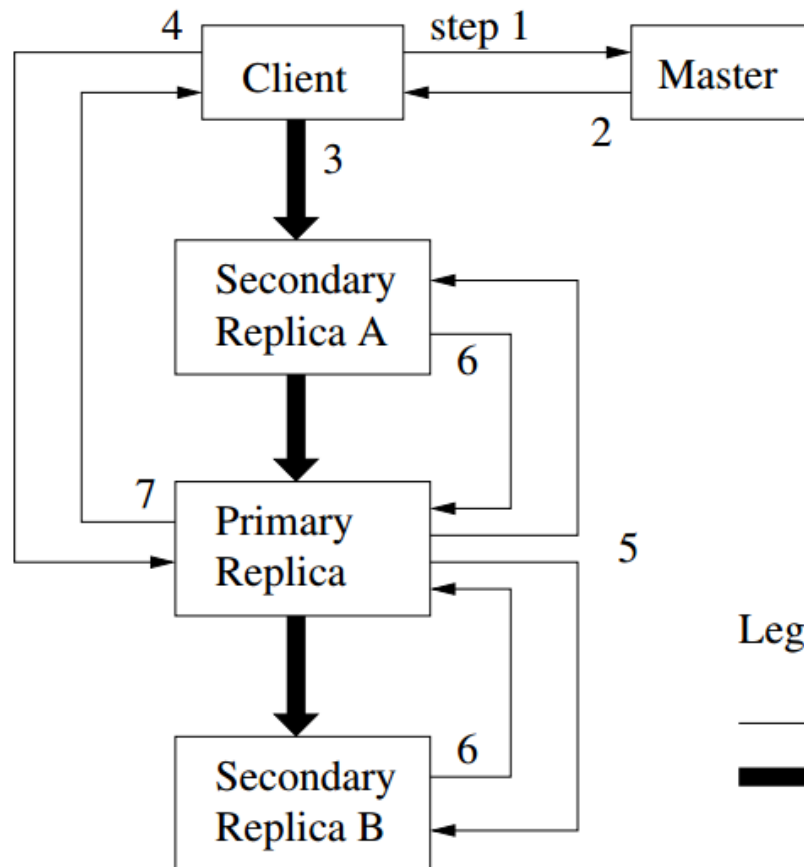


Step3:
push data to
all replicas.
Carefully
picked chain.
Linearly,
Pipelining.
Full duplex

Legend:


→ Control
→ Data

Question 6 – How the components interact



Step4:
After 3 totally
done.
Client sends
write request
to prim.
Prim applies
mutations.

Question 6 – How the components interact

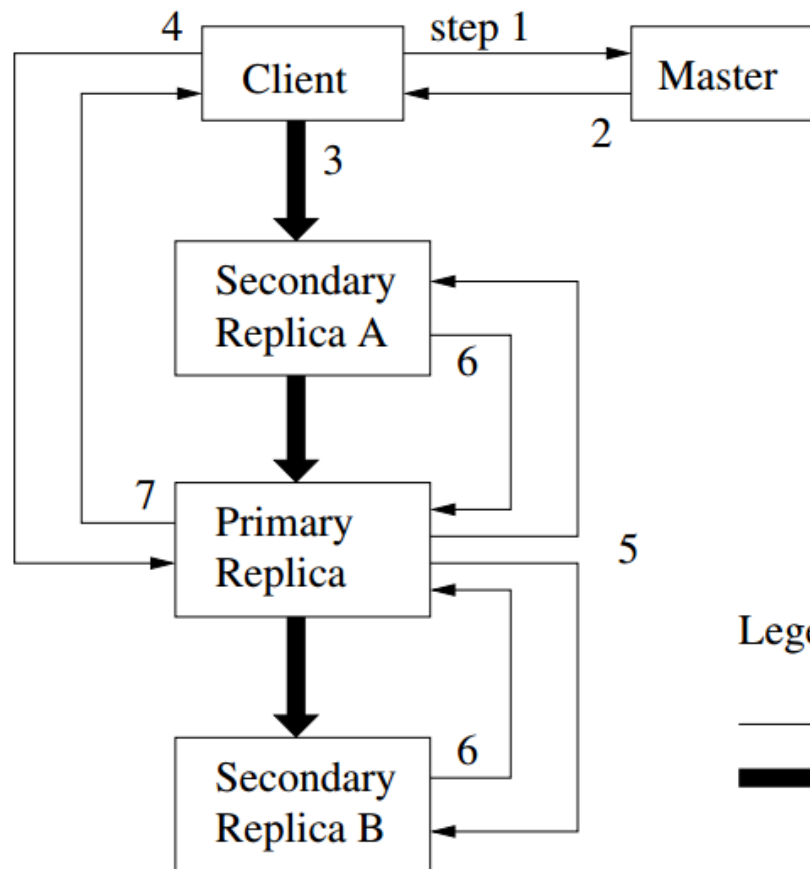


Step5:
Prim forward
the request to
all other
replicas.
Others apply.

Legend:


→ Control
→ Data

Question 6 – How the components interact



**Step6:
Reply to prim**

Question 6 – How the components interact



Step7:
Prim replies to client.
If any error, client retires.
(some secondary replica might fail, inconsistency happens.)

Some detailed questions to discuss

Question 7 –

How is the consistency model?

- **Relaxed consistency model**
- Two types of mutations
 - **Writes**
 - Cause data to be written at an application-specified file offset
 - **Record appends**
 - Operations that append data to a file
 - Cause data to be appended **atomically at least once**
 - Offset chosen by GFS, not by the client

	Write	Record Append
Serial success	<i>defined</i>	<i>defined interspersed with inconsistent</i>
Concurrent successes	<i>consistent but undefined</i>	
Failure	<i>inconsistent</i>	

Question 7.1 – How to append records atomically?

- Be similar to the preceding graph
- But checks if the chunk will exceed 64MB
 - If so, it pads and let the client to append again to a new chunk
 - Limited within $\frac{1}{4}$ size once
- **What if one secondary replica fails?**
 - Client retries. Who failed pads and begin at the same offset.
 - As a result, replicas of the same chunk may contain different data, possibly including duplicates of the same record.
 - But GFS guarantees the data is written at least once and at the same offset.

Question 7.2 – Then how to explain the Table 1?

- States of a file region after a mutation
 - *Consistent*
 - All clients see the same data, regardless which replicas they read from
 - *Defined*
 - *consistent* + all clients see what the mutation writes in its entirety
 - *Undefined*
 - *consistent* + but it may not reflect what any one mutation has written
 - *Inconsistent*
 - Clients see different data at different times

	Write	Record Append
Serial success	<i>defined</i>	<i>defined</i> interspersed with <i>inconsistent</i>
Concurrent successes	<i>consistent</i> but <i>undefined</i>	
Failure	<i>inconsistent</i>	

Question 8 – What about snapshot?

- Goals
 - To quickly create branch copies of huge data sets
 - Or to easily checkpoint the current state
- First revoke leases on chunks involved
- Then write operation will ask master and master will
- Copy-on-write
- locally

Question 9 – How to know which replica is stale

- **Version!**
- Replica version updated when:
 - Master grant a new lease
 - Each mutation is applied

Question 10 – Fault Tolerance and Diagnosis

- Fast Recovery
 - Operation log
 - Checkpointing
- Chunk replication
 - Each chunk is replicated on multiple chunkservers on different racks
- Master replication
 - Operation log and check points are replicated on multiple machines
- Data integrity
 - Checksumming to detect corruption of stored data
 - Each chunkserver independently verifies the integrity
- Diagnostic logs
 - Chunkservers going up and down
 - RPC requests and replies

Background – for measurements

- Two clusters within Google
 - Cluster A: Research & Development
 - Read and analyze data, write result back to cluster
 - Much human interaction
 - Short tasks
 - Cluster B: Production data processing
 - Long tasks with multi-TB data
 - Seldom human interaction

Measurements

- Read rates much higher than write rates
- Both clusters in heavy read activity
- Cluster A supports up to 750MB/read, B: 1300 MB/s
- Master was not a bottle neck

Cluster	A	B
Read rate (last minute)	583 MB/s	380 MB/s
Read rate (last hour)	562 MB/s	384 MB/s
Read rate (since restart)	589 MB/s	49 MB/s
Write rate (last minute)	1 MB/s	101 MB/s
Write rate (last hour)	2 MB/s	117 MB/s
Write rate (since restart)	25 MB/s	13 MB/s
Master ops (last minute)	325 Ops/s	533 Ops/s
Master ops (last hour)	381 Ops/s	518 Ops/s
Master ops (since restart)	202 Ops/s	347 Ops/s

Measurements

- Recovery time (one chunkserver killed)
 - 15,000 chunks containing 600GB are restored in 23.2 minutes (replication rate \cong 400MB/s)
- Recovery time (two chunkserver killed)
 - Each containing 16,000 chunks and 660GB
 - 266 chunks – single replica
 - Higher priority – 2 minutes
 - (Total time not told)

Benefits and Limitations

- Simple design with single master
- Fault tolerance
- Custom designed
- Only viable in a specific environment

Conclusion

- Different than previous file systems
- Satisfies needs of the application
- Fault tolerance

Thanks!

- Welcome to ask and discuss!